PriorityBBR: Investigating Transport-Level Semantic Content Prioritization via TCP Modification

Syed Muhammad Aqdas Rizvi – Lahore University of Management Sciences (LUMS)
25100166@lums.edu.pk

August 6, 2025

Abstract

The demand for improved web performance necessitates further exploration of optimization strategies, particularly for scenarios constrained by network bandwidth or user data limits. This proposal describes the investigation of two distinct, yet synergistic, concepts. **First**, it puts forward a strategy of **Semantic Content Prioritization**, wherein web content elements are explicitly assigned priorities reflecting their importance to the immediate user experience. This enables differentiated delivery schedules, ensuring rapid access to high-priority information (e.g., primary text content) while lower-priority elements (e.g., large images, non-critical scripts) are loaded progressively or adaptively based on network conditions. **Second**, this proposal explores the **implementation of this prioritization strategy directly within the Transmission Control Protocol (TCP)**. This involves embedding priority metadata within TCP segments and modifying transport-layer behavior, potentially leveraging the sophisticated network-awareness capabilities of modern congestion control algorithms like BBR [14], to enforce application-defined priorities during data transmission. While fully acknowledging established protocol layering principles and the critical need for TCP stability, this proposal argues that TCP's documented evolution, shown through the adoption of complex mechanisms like BBR, demonstrates its capacity to adapt when significant performance benefits can be realized. This transport-level approach hypothesizes advantages stemming from tighter integration with TCP's real-time network state awareness compared to purely application-layer solutions. This proposal provides a detailed explanation of the motivation, proposed technical mechanisms, anticipated challenges alongside potential refutations, comparative analysis with existing alternatives, and outlines a rigorous research methodology for evaluating the feasibility and efficacy of this approach.

1 Introduction: Description of Core Concepts

1.1 Semantic Content Prioritization

The correlation between web performance and user engagement is unequivocally established [19]. However, conventional delivery mechanisms often fail to differentiate between content elements based on their contribution to the user's perceived experience, especially under network duress. Consider a standard news article page accessed over a slow mobile connection: the immediate value lies in the headline and article text. Yet, network resources might be equally consumed transmitting bytes for large banner images or third-party analytics scripts, delaying the presentation of the primary content.

This proposal presents a strategy centered on **Semantic Content Prioritization**. By assigning explicit priorities derived from developer annotations, server-side heuristics, or potentially automated analysis,

to distinct content elements (HTML structure, critical CSS, primary text, interactive scripts, different image resolutions, advertisements, etc.), the delivery process can be intelligently orchestrated. The objectives of this strategy are to:

- **Optimize Perceived Performance:** Dramatically reduce the time required for the user to access and interact with the most critical content elements, enhancing subjective speed.

- **Enable Robust Progressive Enhancement:** Facilitate the rapid rendering of a functional baseline page, with subsequent loading of non-essential enhancements contingent on available bandwidth.

- **Facilitate Adaptive Quality Delivery:** Allow for initial delivery of lower-fidelity content (e.g., low-resolution images) with subsequent "upgrades" to higher fidelity versions if network conditions permit.

- **Support Constrained Data Scenarios:** Provide mechanisms to limit delivery to only the highest-priority content tiers, respecting user data caps or explicit low-data preferences.

This strategy fundamentally shifts the focus from undifferentiated byte delivery to a value-driven transmission schedule aligned with user perception.

1.2 Investigating TCP Implementation

The second concept addresses the *implementation venue* for this prioritization strategy. I propose a focused investigation into embedding the enforcement mechanism directly within the **Transmission Control Protocol (TCP)**. The rationale for exploring this seemingly unconventional approach stems from the potential to leverage TCP's unique characteristics:

- **Comprehensive Network State Awareness:** TCP, particularly when employing advanced congestion control algorithms like BBR, maintains a real-time model of network path characteristics, including estimates of bottleneck bandwidth (BtlBw) and round-trip propagation time (RTprop) [14]. Integrating prioritization logic at this layer could enable feedback loops that react more rapidly and accurately to dynamic network conditions than application-level logic relying on potentially delayed or abstracted signals.

- **Direct Control Over Transmission Mechanics:** TCP manages data segmentation, send buffering, transmission scheduling (pacing), and retransmission logic [6, 34]. Implementing prioritization here offers the potential for fine-grained control over precisely which data segment is transmitted next when capacity is limited.

- **Standardization Potential:** Although achieving standardization is a significant hurdle, a successful TCP-level mechanism could offer broader applicability across diverse applications compared to potentially fragmented solutions tied to specific application protocols (e.g., HTTP variants).

This proposal seeks to rigorously assess whether these theoretical advantages manifest as practical benefits sufficient to justify the inherent challenges of modifying a foundational internet protocol. Furthermore, while strongly motivated by enhancing web performance, such a transport-level prioritization mechanism could potentially benefit a wider range of interactive TCP-based applications where data streams inherently contain elements of varying semantic importance for responsiveness.

1.3 Addressing Foundational Objections: Precedent and Philosophy

Modifying TCP to incorporate application-level semantics (content priority) immediately confronts valid architectural and practical objections, which must be addressed proactively:

- **Protocol Layering Principles:** The canonical OSI model emphasizes separation of concerns [42]. Introducing application-specific hints into the transport layer appears to violate this principle. *Response:* While layering is a vital design guideline, performance optimization in real-world systems often necessitates pragmatic cross-layer interactions. TCP's congestion control is already influenced by application sending patterns. QUIC integrates transport and encryption [27]. BBR actively models end-to-end path properties [14]. We argue that passing priority *metadata* for TCP to use in its existing resource management functions (scheduling, buffer management) represents a justifiable, pragmatic evolution, not a fundamental architectural breach, provided substantial benefits are demonstrated.

- **TCP Complexity and Ossification:** There exists a strong and valid imperative to maintain TCP's stability and resist unnecessary complexity, often termed "TCP ossification" [cf. 25]. *Response:* TCP has demonstrably *not* remained static. The development and widespread deployment of complex algorithms like CUBIC [23] and particularly BBR [14] underscore the ecosystem's willingness to adopt significant internal complexity when driven by clear performance imperatives in evolving network environments [cf. 7, 17]. BBR's sophisticated path modeling logic is a testament to TCP's capacity for evolution. This proposal asserts that semantic priority awareness could be the next such performance-driven evolution, following established precedent. Further, research into synthesizing existing transport strategies for optimal performance in demanding environments like data centers, such as PASE [31], also demonstrates the continued effort to evolve and combine transport mechanisms beyond their original singular designs, supporting the idea that transport protocols can and do adapt to new requirements.

- **The Rise of QUIC/HTTP/3:** QUIC was explicitly designed to facilitate transport innovation by operating over UDP and integrating features like stream multiplexing without head-of-line blocking [27], enabling effective application-level prioritization in HTTP/3 [12, 30]. Does this render TCP modification moot? *Response:* The success of QUIC/HTTP/3 validates the *need* for robust prioritization. However, it doesn't automatically preclude potential advantages of a TCP-integrated approach, specifically regarding the immediacy of network state feedback within the TCP stack itself. Furthermore, QUIC currently lacks standardized transport-level prioritization mechanisms applicable beyond HTTP [cf. 26]. Investigating TCP remains relevant for the vast existing TCP ecosystem and for potentially informing future QUIC evolution.

This research proceeds with full awareness of these objections, framing the investigation as an exploration of potential performance trade-offs rather than a premature assertion of TCP's superiority.

1.4 Research Objectives

The primary objectives of this proposed research are:

1. To design and specify concrete mechanisms for embedding semantic priority metadata within TCP

segments and modifying TCP's transmission logic (particularly within a BBR context) to act upon this metadata while incorporating explicit fairness controls.

2. To quantitatively evaluate, through simulation and prototype implementation, the impact of the proposed mechanism on key web performance metrics (e.g., FCP, LCP [22], TTI) under a representative range of network conditions.

3. To rigorously assess the fairness implications of the proposed modifications, both within a single prioritized flow (intra-flow fairness) and between prioritized flows and competing standard TCP/QUIC flows (inter-flow fairness) using established metrics [e.g., 29].

4. To conduct a comparative performance and complexity analysis against state-of-the-art application-layer prioritization techniques, specifically HTTP/3 prioritization over QUIC.

5. To identify and analyze the practical challenges related to implementation, standardization, and deployment.

2 Technical Background

2.1 TCP Fundamentals

The Transmission Control Protocol provides the bedrock for reliable internet communication, offering connection-oriented, ordered, and error-checked byte stream delivery [34]. Its core functions include the three-way handshake, sequence numbering, acknowledgments, checksums for integrity, sliding window flow control, and sophisticated congestion control mechanisms [6] designed to prevent network collapse [28]. Critically for this proposal, standard TCP operates without intrinsic knowledge of the semantic meaning or relative importance of the application data it transports.

2.2 Evolution of TCP Congestion Control

TCP congestion control has evolved significantly. Early algorithms (Tahoe, Reno) were primarily loss-reactive. CUBIC [23] improved performance in high bandwidth-delay product networks but largely retained the loss-based reaction model. TCP BBR [14] marked a fundamental shift to a model-based approach, actively probing the network to estimate BtlBw and RTprop. BBR attempts to control the amount of data in flight to match the estimated BDP, cycling through distinct phases (STARTUP, DRAIN, PROBE_BW, PROBE_RTT) to manage throughput and queuing delay [13]. This inherent network awareness makes BBR a particularly interesting foundation for integrating priority-based decisions.

2.3 Application-Layer Prioritization Solutions

Recognizing the need for prioritization, application protocols have incorporated relevant mechanisms:

- **HTTP/2** introduced stream multiplexing, header compression (HPACK), server push, and stream prioritization based on dependencies and weights [11]. However, its effectiveness is often hampered by inconsistent implementations and, fundamentally, by TCP's head-of-line blocking phenomenon, where the loss of a single TCP segment can stall all multiplexed HTTP/2 streams.

- **HTTP/3 and QUIC** represent a significant advance. By running over QUIC [27], which provides multiple independent, individually flow-controlled transport streams, HTTP/3 [12] largely overcomes the transport-level HoL blocking issue. This allows the application-level prioritization schemes (e.g., as defined in RFC 9218 [30]) to operate much more effectively [17]. QUIC also introduces benefits like reduced connection latency and integrated encryption.

2.4 Limitations of Existing Network-Layer QoS

IP-level mechanisms like the Differentiated Services Code Point (DSCP) field [32] allow packets to be marked for preferential treatment. However, their efficacy for end-to-end application prioritization across the public internet is limited due to inconsistent honoring of these markings by intermediate routers and the fundamental disconnect between network-layer packet handling and transport-layer stream semantics.

2.5 Related Work in Transport-Layer Scheduling

While our primary focus is on enhancing user-perceived performance for general internet traffic, particularly web content, it is instructive to note that transport-level modifications for prioritization and specialized scheduling have been extensively explored in other contexts, notably data center networks. Schemes like pFabric [4] and PIAS [8, 9] have demonstrated notable benefits by incorporating concepts such as flow size awareness, preemption, and information-agnostic scheduling to optimize for metrics like flow completion times. While these systems are tailored for the unique characteristics of datacenter traffic (e.g., many short flows, low RTTs), they establish a strong precedent for the effectiveness of modifying transport layer behavior to achieve specific performance objectives through prioritization. Similarly, PASE [31] explores synthesizing transport strategies for near-optimal performance, indicating a research thrust towards more adaptive and multi-faceted transport protocols. As such, the work herein draws inspiration from this broader willingness to evolve transport mechanisms while targeting a different problem domain and traffic characteristics (semantic importance in user-facing applications over general internet paths).

3 High-Level Design and Architecture

3.1 Priority Definition and Signaling (Application Domain)

The assignment of priorities occurs before data reaches TCP:

- **Assignment:** Priorities (e.g., numerical levels: 0 = Critical, 1 = High, 2 = Normal, 3 = Low, 4 = Background) are assigned by developers (via HTML attributes like fetchpriority [39] or custom data-transport-priority, CSS comments, JavaScript framework integration), server-side logic (based on content type, URL patterns, heuristics), or potentially advanced models (ML/LLM analysis, with performance considerations).

- **API Extension:** A mechanism is required to pass this priority metadata from the application to the TCP stack alongside the data buffers submitted via the socket API (e.g., extending sendmsg() with control messages or defining a new setsockopt() level).

Developer Mechanisms for Priority Assignment

To make the priority assignment concrete, consider these illustrative mechanisms a developer might use:

- **HTML Attributes:** Leveraging existing or custom attributes.

    - Using the standard fetchpriority hint [39] primarily informs the browser's fetch scheduler, but a server could potentially use this hint to infer transport priority:

    ```
    <link rel="stylesheet" href="critical.css"
        fetchpriority="high">
    <img src="hero.jpg" fetchpriority="high">
    <script src="analytics.js"
        fetchpriority="low"></script>
    ```

    - Using custom `data-` attributes to explicitly signal desired transport priority, potentially alongside adaptive loading hints:

    ```
    <img src="high-res.jpg" data-transport-priority="low"
        data-lowres-src="low-res.jpg"
        data-transport-priority-lowres="medium">
    ```

    In this hypothetical example, the server application would parse these attributes and associate the corresponding priority level when sending image data chunks to the TCP socket.

- **Conceptual JavaScript Framework Integration:** Modern frameworks could abstract priority assignment.

    - Components might accept priority props:

    ```
    <Image src="background.png"
        transportPriority={PRIORITY.BACKGROUND} />
    <ArticleText content={...}
        transportPriority={PRIORITY.CRITICAL} />
    ```

    - Data fetching libraries could allow specifying priority:

    ```
    fetchResource('/api/data',
        { transportPriority: PRIORITY.HIGH });
    ```

    These framework-level assignments would translate into setting the appropriate priority via the extended socket API when the underlying data is transmitted.

- **Server-Side Logic:** A backend application or CDN edge worker could determine priority based on request properties or content analysis before writing data to the socket:

    ```
    if (isCriticalApiPath(request.url)) {
      socket.write(data, { transportPriority: PRIORITY.CRITICAL });
    } else {
    ```

```
    socket.write(data, { transportPriority: PRIORITY.NORMAL });
}
```

It is important to distinguish this proposed mechanism from purely application-layer HTTP prioritization schemes like RFC 9218 [30]. While HTTP headers defined in such schemes could *inform* the server application's decision on what priority to signal to the TCP layer via the socket API, the headers themselves are not directly interpreted by TCP. The priority information must be associated with the data buffers passed to the transport layer.

Additionally, this concept of decomposing a larger application-level task (e.g., rendering a webpage) into components of varying importance aligns conceptually with research in other domains, such as coflow scheduling in data centers [15, 16, 20]. In coflow scheduling, multiple related network flows contributing to a single distributed computation are managed collectively to optimize overall job completion. Similarly, a webpage can be viewed as a "coflow" of resources (HTML, CSS, scripts, images), where the timely delivery of critical components dictates the user's perceived performance. This proposal aims to enable finer-grained, segment-level enforcement of such priorities within a single TCP connection carrying these diverse resources.

3.2 Embedding Priority Metadata in TCP Segments

- **Primary Approach: TCP Option.** Define a new TCP Option Kind (requiring IANA assignment). This option would be included in data-carrying TCP segments.
  Format: Must be compact due to the 40-byte total limit for all options [34]. A 2-byte option could suffice: 1 byte for Kind, 1 byte for Length=2, leaving payload bits within the Kind/Length bytes or requiring a subsequent byte for the priority value (e.g., encoding 3-5 priority levels in a few bits).
  Challenges: Option space scarcity; potential for middlebox interference (filtering or stripping unknown options) [25].

- **Alternative Approaches (Less Favored):** Signaling priority during the SYN handshake limits granularity; establishing a dedicated control channel using options adds significant complexity.

The per-segment TCP Option approach offers the most promising balance of granularity and feasibility, despite the known challenges with TCP options.

3.3 Modifying TCP Transmission Logic (Priority Enforcement within BBR Context)

The core innovation lies in modifying TCP's behavior, hypothesized here within a BBR framework:

- **Priority-Aware Send Buffer Management within the Kernel:** The core modification resides within the TCP stack, typically operating within the operating system kernel. When an application writes data to a socket using the extended API (which includes priority metadata), TCP places this data (or references to it along with its associated priority) into its kernel-managed **TCP send buffer**. The proposed priority-aware TCP logic then implements a priority queuing discipline on this kernel buffer. Instead of strict FIFO selection, when the congestion window (cwnd), pacing rate, and receiver window (rwnd) permit transmission, the modified TCP logic preferentially

7

selects data chunks associated with the highest available priority level from the send buffer for segmentation and subsequent transmission. Lower-priority data is only segmented and sent when no higher-priority data is pending in the buffer or as dictated by fairness mechanisms (see below). This ensures scheduling decisions happen as close to the transmission time as possible, informed by both application priority and current network conditions.

- **Integration with BBR Dynamics:**
  Bandwidth Allocation: When BBR estimates BtlBw, the priority scheduler ensures this capacity is preferentially used for segments carrying higher-priority data.
  Congestion Response Modification: Could BBR's reaction to congestion signals (loss, ECN, RTT increase suggesting queue growth) be modulated by the priority of data being sent or ACKed? For example, could the multiplicative decrease factor be slightly less severe if only high-priority data was in flight? This requires rigorous analysis to avoid unfairness.
  PROBE_BW Phase: When BBR probes for additional bandwidth, could it preferentially use higher-priority segments for these probes, ensuring critical data benefits first from newly discovered capacity?
  PROBE_RTT Phase: Ensure that the brief rate reduction during PROBE_RTT doesn't unduly delay critical, high-priority segments waiting in the buffer.

- **Facilitating Adaptive Loading:** This mechanism inherently supports adaptive strategies. The application sends data for a low-resolution image (medium priority). Once ACKed, if BBR indicates sufficient bandwidth, the application sends data for the high-resolution version (low priority), which TCP will transmit only after pending higher-priority data is cleared.

- **Supporting Data Constraints:** An application, informed by user preference, could signal via the extended API that data below a certain priority level should not be transmitted by TCP for this connection. TCP would then refrain from sending segments associated with those lower priorities.

- **Mandatory Fairness Mechanisms:** To prevent starvation and ensure equitable network resource usage:
  Intra-flow Fairness: Implement safeguards ensuring low-priority data eventually progresses (e.g., assigning a small, guaranteed fraction of the sending opportunity, priority aging).
  Inter-flow Fairness: Strictly limit the aggregate advantage conferred by prioritization. The overall flow must remain "TCP-friendly" relative to standard TCP flows (CUBIC, Reno) and fair to competing BBR flows. Congestion control modifications must prioritize stability and fairness over aggressive prioritization [cf. 29].

3.4 Refuting Anticipated Technical Objections

- **Middlebox Interference [25]:** While real, the impact might be mitigated by increasing encryption (limiting deep inspection) and focusing on the tangible benefits achievable between two cooperating, updated endpoints (client and edge server). The primary gain is sender-side scheduling.

- **Compounding BBR Fairness Issues [24, 36]:** Acknowledged. The design *must* incorporate explicit fairness controls as a primary constraint, potentially making the prioritized version *fairer* than baseline BBR under certain conditions by managing internal buffer contention more intelligently. The goal is not to make BBR more aggressive, but more semantically aware in its scheduling.

- **Priority Definition Complexity/Gaming:** This remains an application-level challenge. The TCP mechanism provides the *enforcement* capability; defining the *policy* is external. Initial deployments can use simple, coarse-grained priorities. Mechanisms to prevent gaming (e.g., browser policies) can evolve alongside adoption.

- **Interaction with Application Protocols (HTTP) and Potential Delays:** Concerns may arise about how prioritized TCP interacts with protocols like HTTP, particularly regarding potential delays or "hanging" of requests.

- **Nature of Delay:** It is necessary to understand that delaying lower-priority segments to favor higher-priority ones *within the same TCP connection* under bandwidth constraints is the *intended behavior* of this proposal. This is sender-side scheduling, distinct from traditional receiver-side TCP Head-of-Line (HoL) blocking caused by packet loss.

- **HTTP/1.1:** In sequential HTTP/1.1 requests on a persistent connection, prioritization primarily affects the delivery order of segments within a single large response. Pipelined requests (rarely used effectively) could see a later request delayed if an earlier response contains low-priority data consuming limited TCP sending capacity.

- **HTTP/2 Multiplexing:** HTTP/2 multiplexes multiple requests/responses (streams) over a single TCP connection [11] and has its own stream prioritization logic. The proposed TCP prioritization operates at a layer below HTTP/2 streams. When TCP bandwidth is limited, it will schedule segments based on the priority assigned via the socket API, regardless of which HTTP/2 stream the segment technically belongs to. This creates a potential interaction:

  - **Coordination is Key:** Ideally, the priority assigned at the TCP socket level should directly correspond to the application-level priority (e.g., derived from HTTP/2 stream weights or RFC 9218 priorities). If a resource is high priority in HTTP/2, the server application *must* signal high priority to TCP when sending its data. Consistent signaling ensures smooth coordinated operation.

  - **Potential Conflict:** If priorities conflict (e.g., HTTP/2 signals low priority, but the server mistakenly tells TCP socket the data is high priority), the TCP-level priority would likely dominate the actual transmission order when bandwidth is scarce, potentially undermining the intended HTTP/2 schedule. This serves to emphasize the need for careful implementation in the server application logic that bridges HTTP requests to TCP socket writes.

- **Mitigation via Fairness:** The mandatory intra-flow fairness mechanisms are indispensable here. They ensure that even the lowest-priority data within the connection eventually makes progress, preventing indefinite stalls or "hanging" and guaranteeing eventual completion of all data transmission, albeit potentially delayed.

- **Stateless HTTP:** HTTP's stateless request-response nature is orthogonal to this. TCP provides the stateful connection over which these requests flow. The proposal modifies how data is scheduled *within* that stateful TCP connection.

Therefore, while prioritization intentionally introduces relative delays for lower-priority data, careful coordination between application-level priority signals and socket-level signals, along with robust transport-level fairness mechanisms, is necessary to prevent pathological blocking and ensure predictable behavior, especially when multiplexing protocols like HTTP/2 are used.

3.5 The Essential Role of CDNs and Edge Infrastructure

Content Delivery Networks and Edge computing platforms are natural deployment points:

- They can centralize priority assignment logic (heuristics, processing developer hints).

- They can perform necessary content transformations (e.g., generating multi-resolution image variants linked to priorities.

- They can deploy the modified TCP stack on their edge servers, impacting communication with potentially unmodified clients (who still benefit from prioritized sending).

4 Low-Level Mechanism Design

This section details the preliminary low-level design for implementing PriorityBBR. This specification serves as a concrete foundation for the proposed ns-3 simulation work and future prototyping. The design is intended to be both effective and minimally invasive to the core TCP state machine in its initial form.

4.1 Extended Socket API Specification

To communicate priority from the application to the transport layer, the socket API must be extended. We propose a mechanism inspired by existing ancillary data messages (`sendmsg`/`recvmsg`) and socket options (`setsockopt`).

- **Per-Message Priority (Primary Mechanism):** The most flexible approach is to specify priority on a per-write basis. This will be modeled using a custom Tag that can be attached to packets in ns-3's application layer. In a real system, this would be analogous to using `sendmsg()` with a control message.

  - **Model:** An application preparing data to send would also create a `PriorityTag` containing a priority value.
  - **Priority Levels (Initial):** We propose five distinct priority levels for initial simulation, mapping well to the use cases:
    ```
    0: PRIO_CONTROL (e.g., C2, SSH keystrokes, XR tracking)
    1: PRIO_INTERACTIVE (e.g., critical web content, API responses)
    2: PRIO_NORMAL (e.g., normal images, non-critical data)
    3: PRIO_BULK_ADAPTIVE (e.g., adaptive video, high-res images)
    4: PRIO_SCAVENGER (e.g., background updates, logs)
    ```

- **Per-Socket Default Priority (Fallback):** An application can set a default priority for a socket. Any data written without a per-message priority tag will inherit this default.

  - **Model:** This will be modeled as a new attribute on the ns-3 `TcpSocketBase` class. In a real system, this would be a `setsockopt()` call, e.g., `setsockopt(fd, IPPROTO_TCP, TCP_PRIORITY, &prio_level, sizeof(prio_level));`.

This two-tiered approach provides both flexibility for complex applications and simplicity for basic ones.

4.2 TCP Option Format for Inter-Host Communication

While our initial focus is on sender-side scheduling, a complete design includes a mechanism for signaling priority to the network and the receiver. We define a new TCP Option for this purpose.

- **Kind:** Experimental (e.g., 254, as per RFC 4727).

- **Length:** 3 bytes. This is a minimal, fixed-length option.

- **Info (1 byte payload):** The structure of the Info byte is as follows:

  - **Priority Level (3 bits):** The 3 most significant bits (MSB) encode up to 8 priority levels. This directly maps to our 5 proposed levels.

  - **Reserved (5 bits):** The remaining 5 bits are reserved for future use (e.g., flags for requesting no-drop, etc.).

Example Format: `[Kind=254][Length=3][P2 P1 P0 R4 R3 R2 R1 R0]` This format is compact and aligns with TCP header constraints [34]. While middlebox interference is a known issue [25], defining the option is crucial for a complete architectural proposal.

4.3 Priority-Aware TCP Send Buffer and Scheduling Algorithm

The core logic modification occurs within the TCP send buffer management and the segment transmission scheduling. We will modify the ns-3 `TcpSocketBase` send buffer logic.

- **Buffer Structure:** Instead of a single FIFO buffer, the send buffer will be logically partitioned into multiple queues, one for each priority level. In ns-3, this can be implemented as five distinct `std::list` or `std::deque` objects. When the application sends data with a `PriorityTag`, the data is enqueued into the corresponding priority queue.

- **Scheduling Algorithm (Pseudocode):** The main TCP sending logic (often called 'SendPacket', 'SendPendingData', or similar in a TCP implementation) is modified. This function is called when the TCP state machine determines it is eligible to send new data (i.e., the send window is open, pacing allows it, etc.).

```
function schedule_and_send():
  // Pre-condition: BBR logic has determined a "send budget":
  //   can_send_bytes = min(cwnd - in_flight, pacing_allowance)

  while can_send_bytes > 0:
    segment_to_send = NULL

    // 1. Check for pending retransmissions (highest implicit priority)
    if retransmit_queue is not empty:
      segment_to_send = retransmit_queue.pop()
    else:
      // 2. Iterate through priority queues from HIGHEST to LOWEST
      for prio_level from PRIO_CONTROL to PRIO_SCAVENGER:
```

```
        if priority_queue[prio_level] has data:

          // 3. Apply Intra-flow Fairness (Starvation Prevention)
          if prio_level > PRIO_INTERACTIVE: // For non-critical data
            if not fairness_allows_tx(prio_level):
              continue // Skip this level for now

          segment_to_send = create_segment_from(priority_queue[prio_level])
          break // Found a segment to send

    if segment_to_send is NULL:
      break // No eligible data to send

    // 4. Transmit the segment
    transmit(segment_to_send)
    can_send_bytes -= segment_to_send.size

    // 5. Update fairness state
    update_fairness_state(segment_to_send.priority)
```

- **Intra-flow Fairness Mechanism (`fairness_allows_tx`):** To prevent starvation of low-priority data, a simple but effective mechanism is essential for the initial simulation. We propose a credit-based or token bucket scheme per priority level.

  - Each time a high-priority segment is sent, a small "debt" is accrued for the lower-priority queues.

  - A low-priority queue can only be serviced if it has data AND its "debt" is below a certain threshold (or it has accrued enough "credits" from periods of inactivity).

  - A simpler alternative is a timer-based override: if the data at the head of a low-priority queue has been waiting for more than a threshold time (e.g., 500ms), it is temporarily promoted to a higher priority for scheduling. This will be the initial approach for the ns-3 strawman.

4.4   Interaction with BBR State Machine

For the initial ns-3 implementation, PriorityBBR will treat the BBR algorithm as a "black box" that provides a sending budget.

- **Phase 1 (Scheduling Only):** Our modified scheduler (from the pseudocode above) operates within the constraints calculated by the unmodified BBRv1 or BBRv2 implementation in ns-3. BBR will determine the pacing rate and congestion window (`cwnd`); our scheduler then uses that budget to send the highest-priority available data. This allows us to isolate and evaluate the effects of scheduling alone.

- **Phase 2 (Future Research - State Machine Interaction):** Subsequent research, as part of the advanced simulation scenarios, will investigate modifying BBR's internal logic. This involves exploring questions such as:

– Should BBR use high-priority packets when probing for bandwidth in the 'PROBE_BW' phase to ensure that any newly discovered capacity is immediately utilized by critical data?

– How should congestion events (e.g., loss of a high-priority vs. low-priority packet) influence BBR's estimation of `BtlBw` or its decision to exit the STARTUP phase?

This phased approach ensures a methodical investigation, starting with the simplest effective change.

5 Illustrative Use Cases

While enhancing web performance is a primary motivator for this research, the fundamental concept of transport-level semantic prioritization holds significant potential across diverse domains. The use cases can be categorized based on their operating environment and objectives, ranging from mission-critical tactical networks to everyday interactive applications.

5.1 High-Stakes Environments: Defense and Tactical Networks

The proposal's core concept finds a compelling and powerful application in defense and tactical communication environments. Drawing inspiration from the Internet's origins in resilient, packet-switched defense networks, this work addresses the critical challenges of modern tactical networks. These environments, often characterized as Mobile Ad-hoc Networks (MANETs), operate with constrained, lossy, and variable wireless links, yet must reliably carry heterogeneous traffic of mixed criticality [2].

PriorityBBR can be designed to enforce mission-critical precedence at the transport layer, directly on the sending device, before data even enters the wider network.

- **Command and Control (C2) Precedence:** Tiny, latency-sensitive C2 messages (e.g., targeting updates, threat alerts, remote control signals for UAVs) can be assigned the highest priority. This ensures they are transmitted from the end-system's buffer before pending, larger data chunks from lower-priority streams, such as high-resolution video feeds or bulk logistical data.

- **Interaction with In-Network QoS:** This end-host scheduling mechanism is a powerful complement to existing in-network prioritization architectures like the IP Differentiated Services (DiffServ) framework [18], which forms the basis for military Multi-Level Precedence and Preemption (MLPP). The complete workflow could be:

  1. **PriorityBBR** ensures a critical segment is sent from the host first.
  2. The IP layer marks that same packet with a high-precedence DSCP value.
  3. Routers in the tactical network give the packet preferential treatment based on its DSCP mark.

  This creates a true end-to-end precedence framework, from the application buffer to the final destination.

- **Enhanced Situational Awareness:** By enabling the prioritization of Blue Force Tracking updates or critical sensor alerts over less timely data, the mechanism directly contributes to a more accurate and real-time common operational picture. This is a cornerstone of modern military doctrine such as Network-Centric Warfare (NCW) [3].

5.2 Latency-Sensitive Interactive Applications

Beyond defense, the principle of preserving interactivity under load is critical for many established remote access and real-time collaboration tools that rely on TCP.

- **Remote Desktop Protocols (e.g., RDP, VNC over TCP):** User input events (keystrokes, mouse movements) are extremely latency-sensitive and demand the highest priority. Screen updates directly resulting from those actions can be prioritized next, while updates to static screen regions or background file transfers are assigned lower priorities to prevent interactive lag.

- **Interactive Shells (SSH):** For protocols like SSH [41], echoing user keystrokes and immediate command responses are critical for interactivity. Conversely, bulk data transfers over the same connection (e.g., via SFTP or displaying a large file) can be marked with lower priority to keep the shell responsive.

- **Real-time Collaboration and Messaging Systems:** In collaborative editing tools, real-time text updates can be prioritized over presence notifications, synchronization of non-visible document sections, or large embedded media file transfers.

- **Other Interactive Data Streams:** Similar benefits apply to interactive database queries (prioritizing initial results over large data set fetches) or financial data feeds (prioritizing critical market updates over auxiliary data).

5.3 Future Applications in 5G/6G Systems

The architectural shift in 5G and the forward-looking vision for 6G are defined by the need to support a diverse set of services with conflicting performance requirements on a unified infrastructure [1]. The co-existence of Enhanced Mobile Broadband (eMBB), massive Machine-Type Communications (mMTC), and Ultra-Reliable Low-Latency Communication (URLLC) creates a fundamental resource management challenge [33]. This end-host prioritization becomes particularly effective with two core 5G/6G architectural enablers, which are Network Slicing and Multi-access Edge Computing (MEC) [37]. PriorityBBR acts as the crucial "first-mile" scheduler, in that it ensures that an application's most latency-sensitive data is the first to be submitted to its designated high-performance network slice, maximizing the end-to-end performance benefits of these new architectures.

- **Extended Reality (XR) for Mobile and Aerial Platforms:** XR applications are a primary driver for 5G and future networks, imposing stringent demands for both high bandwidth and ultra-low latency. These challenges are amplified when delivering XR experiences via mobile platforms like Unmanned Aerial Vehicles (UAVs) for applications in public safety, industry, or defense [10]. An XR application running on such a platform generates a mixed-criticality data stream:
  - **Highest Priority (URLLC-like):** Device pose (head/controller tracking), user input, and real-time control signals. These are typically small, but their delivery is extremely latency-sensitive to prevent motion sickness and ensure accurate operation.
  - **Lower Priority (eMBB-like):** High-resolution textures, complex 3D models streamed from the edge, environment maps, and other non-interactive assets. These are often large but can tolerate higher latency.

Without transport-level prioritization, the transmission of a large 3D model could block a small but critical tracking update, causing a noticeable stutter that breaks the immersive experience or compromises remote control. PriorityBBR would allow the XR application to assign the highest priority to the latency-sensitive pose data, ensuring it preempts any pending asset data in the TCP send buffer and is dispatched immediately.

- **The Tactile Internet and Remote Control:** The vision of the Tactile Internet, a cornerstone of 6G research, involves transmitting touch and actuation in real-time to control remote robotics or interact with haptic interfaces [21, 35]. This requires end-to-end latencies on the order of 1-10 milliseconds. In such a system, PriorityBBR would provide the mechanism for the application to guarantee that control and haptic signals—the essence of the "tactile" experience—are never delayed by pending video or log data at the transport layer, providing a critical component for achieving the required end-to-end performance.

5.4 Primary Motivating Use Case: Web Performance

Finally, the primary scenarios motivating this research stem from the need to improve everyday web performance, particularly under constrained network conditions.

- **Constrained News Access:** The user, such as one on a 3G connection, sees the article text load within seconds, while a low-resolution lead image appears shortly after. Web fonts, comments, and high-resolution imagery populate gradually, rather than contending equally for initial bandwidth.

- **Mobile E-commerce:** Critical path elements (product title, price, buy button) render rapidly, allowing a user to decide to purchase before all image thumbnails or related product carousels have fully loaded.

- **Image Galleries:** Low-resolution previews for all images load quickly (medium priority). Full-resolution versions (low priority) download in the background or upon user interaction.

- **Low-Data Mode:** A user enables OS-level data saving. The news site, via priority-aware TCP, transmits only the core text and minimal CSS (highest priority), actively preventing transmission of lower-priority image or script data and resulting in significant data savings.

6 Analysis of Challenges and Mitigation Strategies

Successful realization requires overcoming significant hurdles:

- **Backward Compatibility:** Ensuring seamless operation with legacy TCP stacks is non-negotiable. *Mitigation:* Adherence to RFC 793 regarding unknown options; negotiation of capabilities during handshake; initial deployment focusing on server-side benefits.

- **Standardization Effort:** Achieving IETF consensus (e.g., within TCPM WG or a new WG) is a lengthy, demanding process requiring robust technical justification and community support. *Mitigation:* Rigorous research findings, prototype data, clear performance/fairness analysis, addressing layering concerns proactively, potentially starting with Experimental RFCs.

- **Ecosystem Adoption:** Requires updates across diverse OS vendors, CDN providers, potentially browser/application developers.
  Mitigation: Demonstrating compelling benefits; targeting CDNs first; leveraging open-source implementations; providing clear APIs and documentation.

- **End-to-End QoS Limitations:** Lack of guaranteed priority handling by intermediate routers.
  Mitigation: Focus mechanism on end-host send buffer scheduling, providing benefits independent of transit network behavior.

- **Fairness Guarantees:** Ensuring fairness is particularly critical when modifying complex, model-based congestion control algorithms like BBR, whose own fairness characteristics have been a subject of study [24, 36]. The introduction of an additional prioritization dimension necessitates even more rigorous fairness controls and validation, drawing lessons from prior work on specialized transport protocols [e.g., 5, 31] which also had to consider resource allocation among different classes of traffic or objectives.
  Mitigation: Incorporate explicit fairness algorithms (e.g., weighted fair queuing principles, rate limits on priority advantage) into the core TCP modification design; validate extensively via simulation.

- **Priority Policy Definition:** Complexity and potential for misuse in assigning priorities.
  Mitigation: Separate mechanism from policy; promote best practices; potential for client-side policy enforcement (browsers); start with simple, demonstrable use cases.

- **Security Vulnerabilities:** Potential for new attack vectors.
  Mitigation: Thorough security review throughout design and standardization; extensive fuzzing and penetration testing of implementations.

- **Socio-Economic and Stakeholder Resistance:** De-prioritization of certain content types (e.g., advertisements, third-party trackers) could face significant resistance from stakeholders reliant on immediate content display for revenue (e.g., the web advertisement industry) or functionality. This is a non-technical hurdle that can impede standardization and adoption.
  Mitigation/Consideration: Acknowledge this challenge; focus initial investigations on user-centric benefits (faster core content, data savings); frame prioritization as a mechanism that could be configured by applications/users based on policy, rather than inherently penalizing specific content types; Explore models where essential ad framework components might receive higher priority than the ad creatives themselves; standardization efforts would need to involve diverse stakeholders to find acceptable compromises or demonstrate overwhelming user benefit; Further research could quantify the impact of "core content first" on overall user engagement, which might indirectly benefit even advertising if users are less likely to abandon pages.

7   Comparative Analysis with Alternative Prioritization Techniques

The proposed TCP mechanism must be evaluated against existing solutions:

- **HTTP/2 Prioritization [11]:** Limited by TCP HoL blocking and implementation variance.

- **HTTP/3 over QUIC [12, 27, 30]:** Highly effective due to QUIC streams removing transport HoL blocking; current state-of-the-art for application-level web prioritization. The key differentiator

for the proposed TCP approach would be demonstrating advantages derived from tighter integration with transport-layer network state estimation (BBR).

- **Browser Hints (fetchpriority, preload) [39, 40]:** Influence resource discovery and browser scheduling queues, but not TCP-level transmission dynamics for data already in the send buffer. Complementary, not substitutive.

- **Service Workers [38]:** Powerful for caching and request interception, enabling sophisticated application logic, but not designed for fine-grained control over TCP segment scheduling during transmission.

- **CDN Optimizations:** Focus on content transformation and caching at the edge; generally operate without detailed, real-time insight into the client's specific TCP connection dynamics.

The central research question is whether the hypothesized benefit of TCP's direct network awareness translates into measurable performance gains over the highly effective QUIC/HTTP/3 approach, sufficient to warrant the costs of TCP modification.

8 Proposed Research Methodology

A multi-faceted approach is required:

1. **Formal Mechanism Design:** Specify the TCP Option, socket API extensions, priority queue logic for the send buffer, and precise modifications to BBR's state machine and control loops, including fairness algorithms.

2. **Simulation Environment Setup:** The primary evaluation will be conducted using a discrete-event network simulator, likely ns-3, due to its comprehensive TCP/IP modeling capabilities and flexibility.

3. **Initial Simulation Scenarios (Simple Topologies):**

   - Dumbbell topology: To analyze core prioritization behavior, interaction with congestion control (modified BBR), and impact on single-flow metrics under varying bottleneck bandwidth, RTT, and buffer sizes.
   - Parking lot topology: To assess inter-flow fairness between PriorityBBR flows and standard TCP (CUBIC, BBR) flows, and among multiple PriorityBBR flows with different priority mixes.

4. **Advanced Simulation Scenarios (Representative Topologies):** Subsequent simulations will utilize more complex topologies reflecting realistic web access patterns and potentially mobile network characteristics to evaluate performance under more diverse conditions.

5. **Rigorous Fairness Analysis:** Design specific simulation scenarios and testbed experiments aimed at stressing fairness conditions (e.g., high contention, mix of flow types, persistent low-priority data) to validate the effectiveness of embedded fairness mechanisms.

6. **Performance Evaluation:** Perform direct performance and fairness comparisons against QUIC/HTTP/3 implementations under identical workloads and network conditions. Also consider, where

conceptually applicable, the performance principles demonstrated by specialized datacenter transports like pFabric [5] or PIAS [8] as points of reference for what can be achieved with aggressive, context-aware transport scheduling, even if their specific mechanisms are not directly transferable to general Internet paths, to provide aspirational benchmarks for segment-level scheduling effectiveness.

7. **Complexity and Feasibility Assessment:** Analyze the implementation complexity and potential hurdles to standardization and deployment based on design and prototyping experiences.

9   Conclusion and Future Outlook

This proposal argues for a structured investigation into the potential of implementing **Semantic Content Prioritization** at the **TCP transport layer**. While the strategy of prioritization itself holds clear value for improving perceived web performance, the choice of TCP as the implementation venue requires careful justification and rigorous evaluation.

The primary hypothesis is that leveraging TCP's direct access to real-time network state, particularly within the context of advanced congestion control like BBR, may enable more responsive and efficient prioritization enforcement than purely application-layer approaches. The historical evolution of TCP, embracing complexity like BBR for performance gains, provides a precedent for considering such modifications.

However, the challenges associated with altering TCP (compatibility, standardization friction, ecosystem adoption, ensuring fairness, and managing complexity) are substantial and must be central to the investigation. This research aims not to prematurely champion TCP modification, but to scientifically assess its potential benefits and drawbacks relative to state-of-the-art alternatives, most notably QUIC/HTTP/3.

Through detailed design, simulation, prototyping, and rigorous analysis focused on both performance and fairness, this work seeks to provide data-driven insights into the feasibility and desirability of priority-aware TCP. The findings will either contribute evidence supporting a new direction for TCP evolution or reinforce the advantages of handling prioritization at higher layers, potentially informing future developments in both TCP and QUIC ecosystems. Moreover, the investigation holds potential relevance beyond the HTTP ecosystem, informing how other critical interactive TCP applications could achieve better responsiveness and efficiency under diverse network conditions. Ultimately, this investigation contributes to the broader goal of creating a more performant, efficient, and user-responsive internet infrastructure.

References

[1] 3rd Generation Partnership Project (3GPP). System Architecture for the 5G System (5GS). Technical Specification TS 23.501 V19.4.0, 3GPP, jun 2025. URL `https://portal.3gpp.org/desktopmodules/Specifications/SpecificationDetails.aspx?specificationId=3144`. Release 19.

[2] I.F. Akyildiz, Weilian Su, Y. Sankarasubramaniam, and E. Cayirci. A survey on sensor networks. *IEEE Communications Magazine*, 40(8):102–114, 2002. doi: 10.1109/MCOM.2002.1024422.

[3] David S. Alberts, John J. Gartska, Richard E. Hayes, and David A. Signori. *Understanding Information Age Warfare*. CCRP Publication Series, Washington, D.C., 2001. ISBN 1-893723-04-6.

[4] Mohammad Alizadeh, Shuang Yang, Milad Sharif, Sachin Katti, Nick McKeown, Balaji Prab-hakar, and Scott Shenker. pfabric: minimal near-optimal datacenter transport. In *Proceedings of the ACM SIGCOMM 2013 Conference on SIGCOMM*, SIGCOMM '13, page 435–446, New York, NY, USA, 2013. Association for Computing Machinery. ISBN 9781450320566. doi: 10.1145/2486001.2486031. URL https://doi.org/10.1145/2486001.2486031.

[5] Mohammad Alizadeh, Shuang Yang, Milad Sharif, Sachin Katti, Nick McKeown, Balaji Prab-hakar, and Scott Shenker. pfabric: minimal near-optimal datacenter transport. *SIGCOMM Comput. Commun. Rev.*, 43(4):435–446, August 2013. ISSN 0146-4833. doi: 10.1145/2534169.2486031. URL https://doi.org/10.1145/2534169.2486031.

[6] Mark Allman, Vern Paxson, and Ethan Blanton. TCP Congestion Control. RFC 5681, IETF, September 2009. URL https://www.rfc-editor.org/info/rfc5681.

[7] Amazon Web Services, Inc. Tcp bbr congestion control with amazon cloudfront | networking & content delivery, July 2019. URL https://aws.amazon.com/blogs/networking-and-content-delivery/tcp-bbr-congestion-control-with-amazon-cloudfront/.

[8] Wei Bai, Li Chen, Kai Chen, Dongsu Han, Chen Tian, and Weicheng Sun. Pias: Practical information-agnostic flow scheduling for data center networks. In *Proceedings of the 13th ACM Workshop on Hot Topics in Networks*, HotNets-XIII, page 1–7, New York, NY, USA, 2014. Association for Computing Machinery. ISBN 9781450332569. doi: 10.1145/2670518.2673871. URL https://doi.org/10.1145/2670518.2673871.

[9] Wei Bai, Li Chen, Kai Chen, Dongsu Han, Chen Tian, and Hao Wang. Information-agnostic flow scheduling for commodity data centers. In *Proceedings of the 12th USENIX Conference on Networked Systems Design and Implementation*, NSDI'15, page 455–468, USA, 2015. USENIX Association. ISBN 9781931971218.

[10] Mohammed Khaled Banafaa, Ömer Pepeoğlu, Ibraheem Shayea, Abdulraqeb Alhammadi, Zaid Ahmed Shamsan, Muneef A. Razaz, Majid Alsagabi, and Sulaiman Al-Sowayan. A com-prehensive survey on 5g-and-beyond networks with uavs: Applications, emerging technologies, regulatory aspects, research trends and challenges. *IEEE Access*, 12:7786–7826, 2024. doi: 10.1109/ACCESS.2023.3349208.

[11] M. Belshe, R. Peon, and M. Thomson. Hypertext Transfer Protocol Version 2 (HTTP/2). RFC 7540, IETF, May 2015. URL https://www.rfc-editor.org/info/rfc7540.

[12] M. Bishop. HTTP/3. RFC 9114, IETF, June 2022. URL https://www.rfc-editor.org/info/rfc9114.

[13] Neal Cardwell, Yuchung Cheng, C. Stephen Gunn, Soheil Hassas Yeganeh, and Van Jacobson. Bbr: Congestion-based congestion control. *ACM Queue*, 14, September-October:20 – 53, 2016. URL http://queue.acm.org/detail.cfm?id=3022184.

[14] Neal Cardwell, Yuchung Cheng, C. Stephen Gunn, Soheil Hassas Yeganeh, and Van Jacobson. Bbr: congestion-based congestion control. *Commun. ACM*, 60(2):58–66, January 2017. ISSN 0001-0782. doi: 10.1145/3009824. URL https://doi.org/10.1145/3009824.

[15] Mosharaf Chowdhury and Ion Stoica. Efficient coflow scheduling without prior knowledge. *SIGCOMM Comput. Commun. Rev.*, 45(4):393–406, August 2015. ISSN 0146-4833. doi: 10.1145/2829988.2787480. URL https://doi.org/10.1145/2829988.2787480.

[16] Mosharaf Chowdhury, Matei Zaharia, Justin Ma, Michael I. Jordan, and Ion Stoica. Managing data transfers in computer clusters with orchestra. In *Proceedings of the ACM SIGCOMM 2011 Conference*, SIGCOMM '11, page 98–109, New York, NY, USA, 2011. Association for Computing Machinery. ISBN 9781450307970. doi: 10.1145/2018436.2018448. URL https://doi.org/10.1145/2018436.2018448.

[17] Cloudflare. Introducing HTTP/3 Prioritization. Cloudflare Blog, January 2023. URL https://blog.cloudflare.com/better-http-3-prioritization-for-a-faster-web/.

[18] David L. Black and Zheng Wang and Mark A. Carlson and Walter Weiss and Elwyn B. Davies and Steven L. Blake. An Architecture for Differentiated Services. RFC 2475, December 1998. URL https://www.rfc-editor.org/info/rfc2475.

[19] Deloitte Digital / Google. Milliseconds Make Millions. Industry Report, 2020. URL https://www.thinkwithgoogle.com/_qs/documents/9757/Milliseconds_Make_Millions_report_hQYAbZJ.pdf.

[20] Fahad R. Dogar, Thomas Karagiannis, Hitesh Ballani, and Antony Rowstron. Decentralized task-aware scheduling for data center networks. *SIGCOMM Comput. Commun. Rev.*, 44(4):431–442, August 2014. ISSN 0146-4833. doi: 10.1145/2740070.2626322. URL https://doi.org/10.1145/2740070.2626322.

[21] Gerhard P. Fettweis. The Tactile Internet: Applications and Challenges. *IEEE Vehicular Technology Magazine*, 9(1):64–70, 2014. doi: 10.1109/MVT.2013.2295069.

[22] Google. Largest Contentful Paint (LCP). web.dev Documentation, Ongoing. URL https://web.dev/articles/lcp.

[23] Sangtae Ha, Injong Rhee, and Lisong Xu. Cubic: a new tcp-friendly high-speed tcp variant. *SIGOPS Oper. Syst. Rev.*, 42(5):64–74, July 2008. ISSN 0163-5980. doi: 10.1145/1400097.1400105. URL https://doi.org/10.1145/1400097.1400105.

[24] Marcel Hock, Roland Bless, and Martina Zitterbart. TCP BBR in the Wild: Large-Scale Measurement and Analysis. In *Proceedings of the Internet Measurement Conference (IMC '19)*, pages 262–275, 2019. doi: 10.1145/3355369.3355573.

[25] Michio Honda, Yoshifumi Nishida, Costin Raiciu, Adam Greenhalgh, Mark Handley, and Hideyuki Tokuda. Is it still possible to extend tcp? In *Proceedings of the 2011 ACM SIGCOMM Conference on Internet Measurement Conference*, IMC '11, page 181–194, New York, NY, USA, 2011. Association for Computing Machinery. ISBN 9781450310130. doi: 10.1145/2068816.2068834. URL https://doi.org/10.1145/2068816.2068834.

[26] IETF QUIC Working Group. Priority in QUIC Transport · Issue #104 · quicwg/base-drafts. GitHub Issue Discussion, 2016-2017. URL https://github.com/quicwg/base-drafts/issues/104.

[27] J. Iyengar and M. Thomson. QUIC: A UDP-Based Multiplexed and Secure Transport. RFC 9000, IETF, May 2021. URL https://www.rfc-editor.org/info/rfc9000.

[28] V. Jacobson. Congestion avoidance and control. In *Symposium Proceedings on Communications Architectures and Protocols*, SIGCOMM '88, page 314–329, New York, NY, USA, 1988. Association for Computing Machinery. ISBN 0897912799. doi: 10.1145/52324.52356. URL https://doi.org/10.1145/52324.52356.

[29] Raj K. Jain, Dah-Ming W. Chiu, and William R. Hawe. A Quantitative Measure of Fairness and Discrimination for Resource Allocation in Shared Computer Systems. Technical report, Digital Equipment Corporation, September 1984. URL https://www.cs.wustl.edu/~jain/papers/ftp/fairness.pdf.

[30] M. Kühlewind and S. M. B. Priorities. Extensible Prioritization Scheme for HTTP. RFC 9218, IETF, June 2022. URL https://www.rfc-editor.org/info/rfc9218.

[31] Ali Munir, Ghufran Baig, Syed Mohammad Irteza, Ihsan Ayyub Qazi, Alex X. Liu, and Fahad Rafique Dogar. Pase: Synthesizing existing transport strategies for near-optimal data center transport. *IEEE/ACM Transactions on Networking*, 25(1):320–334, 2017. doi: 10.1109/TNET.2016.2586508.

[32] K. Nichols, S. Blake, F. Baker, and D. Black. Definition of the Differentiated Services Field (DS Field) in the IPv4 and IPv6 Headers. RFC 2474, IETF, December 1998. URL https://www.rfc-editor.org/info/rfc2474.

[33] Parvez, Imtiaz and Rahmati, Ali and Guvenc, Ismail and Sarwat, Arif I. and Dai, Huaiyu. A Survey on Low Latency Towards 5G: RAN, Core Network and Caching Solutions. *Commun. Surveys Tuts.*, 20(4):3098–3130, October 2018. ISSN 1553-877X. doi: 10.1109/COMST.2018.2841349. URL https://doi.org/10.1109/COMST.2018.2841349.

[34] Jon Postel. Transmission Control Protocol. RFC 793, IETF, September 1981. URL https://www.rfc-editor.org/info/rfc793.

[35] Walid Saad, Mehdi Bennis, and Mingzhe Chen. A Vision of 6G Wireless Systems: Applications, Trends, Technologies, and Open Research Problems. *IEEE Network*, 34(3):134–142, 2020. doi: 10.1109/MNET.001.1900287.

[36] Kanon Sasaki, Masato Hanai, Kouto Miyazawa, Aki Kobayashi, Naoki Oda, and Saneyasu Yamaguchi. TCP Fairness Among Modern TCP Congestion Control Algorithms Including TCP BBR. In *2018 IEEE 7th International Conference on Cloud Networking (CloudNet)*, pages 1–4, 2018. doi: 10.1109/CloudNet.2018.8549505.

[37] Taleb, Tarik and Samdanis, Konstantinos and Mada, Badr and Flinck, Hannu and Dutta, Sunny and Sabella, Dario. On Multi-Access Edge Computing: A Survey of the Emerging 5G Network Edge Cloud Architecture and Orchestration. *IEEE Communications Surveys & Tutorials*, 19(3):1657–1681, 2017. doi: 10.1109/COMST.2017.2705720.

[38] W3C Service Worker Working Group. Service Workers 1. W3C Recommendation, December 2022. URL https://www.w3.org/TR/service-workers-1/.

[39] W3C Web Incubator Community Group (WICG) / WHATWG. Priority Hints. Specification Draft, Ongoing. URL https://wicg.github.io/priority-hints/.

[40] W3C Web Performance Working Group. Resource Hints. W3C Recommendation, March 2023. URL `https://www.w3.org/TR/resource-hints/`.

[41] T. Ylonen and C. Lonvick. The Secure Shell (SSH) Protocol Architecture. RFC 4251, IETF, January 2006. URL `https://www.rfc-editor.org/info/rfc4251`.

[42] H. Zimmermann. OSI Reference Model - The ISO Model of Architecture for Open Systems Interconnection. *IEEE Transactions on Communications*, 28(4):425–432, 1980. doi: 10.1109/TCOM.1980.1094702.